



**Model1:****Number of Parameters, MAC Operations and Output Size for Model 1**

Layer Name	Output Size	Parameters	MAC Operations
Conv2D	16x16x32	896	295,014,912 (16x16x32x3x3x3x32)
BatchNormalization	16x16x32	256	0
Conv2D	8x8x64	18,496	118,111,488 (8x8x64x3x3x32x64)
BatchNormalization	8x8x64	512	0
Conv2D	4x4x128	73,856	590,252,544 (4x4x128x3x3x64x128)
BatchNormalization	4x4x128	512	0
Conv2D	4x4x128	1,47,584	590,252,544 (4x4x128x3x3x128x128)
BatchNormalization	4x4x128	512	0
Conv2D	4x4x128	1,47,584	590,252,544 (4x4x128x3x3x128x128)
BatchNormalization	4x4x128	512	0
Conv2D	4x4x128	1,47,584	590,252,544 (4x4x128x3x3x128x128)
BatchNormalization	4x4x128	512	0
MaxPooling2D	1x1x128	0	0
Flatten	128	0	0
Dense	128	16,512	16,512
BatchNormalization	128	512	0
Dense	10	1,290	1,280

**Formula for Calculating MACs** (output\_height \* output\_width \* output\_channels) \* (kernel\_height \* kernel\_width \* input\_channels) \* output\_channels

**Did you observe any overfitting? Should the model train for longer, shorter, or about that number of epochs.**

The Validation accuracy is lower than the Training accuracy. The Validation loss is higher than the Training loss. Hence, Model1 is Overfitting.



**Does it correctly label the picture?**

Sample image of Dog is tested on Model 1 and it correctly classified

```
uploaded = files.upload()
img = Image.open("dog.jpg") # Loading Image
img = img.resize((32, 32)) # Resizing
x = np.array(img) # Convert the image to a numpy array
x = np.expand_dims(x, axis=0)
model = keras.models.load_model("model1.h5")
# Make a prediction on the image
pred = model.predict(x)
# Get the predicted class label
class_index = np.argmax(pred)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
class_name = class_names[class_index]
print(class_name)
```

Choose Files dog.jpg

- **dog.jpg**(image/jpeg) - 884038 bytes, last modified: 2/7/2024 - 100% done

Saving dog.jpg to dog.jpg

1/1 [=====] - 0s 192ms/step



dog

## Model2

Layer	Output Shape	Parameters	MAC Operations
Conv2D	(16, 16, 32)	896	294,912 (16x16x32x3x3x3x32)
BatchNormalization	(16, 16, 32)	128	0
DepthwiseConv2D	(8, 8, 32)	288	57,600 (8x8x32x3x3x1)
Conv2D	(8, 8, 64)	2,112	36,864 (8x8x64x1x1x32)
BatchNormalization	(8, 8, 64)	256	0
DepthwiseConv2D	(4, 4, 64)	576	14,400 (4x4x64x3x3x1)
Conv2D	(4, 4, 128)	8,320	9,216 (4x4x128x1x1x64)
BatchNormalization	(4, 4, 128)	512	0
DepthwiseConv2D	(2, 2, 128)	1,152	3,600 (2x2x128x3x3x1)
Conv2D	(2, 2, 128)	16,512	512 (2x2x128x1x1x128)
BatchNormalization	(2, 2, 128)	512	0
DepthwiseConv2D	(1, 1, 128)	1,152	3,600 (1x1x128x3x3x1)
Conv2D	(1, 1, 128)	16,512	128 (1x1x128x1x1x128)
BatchNormalization	(1, 1, 128)	512	0
DepthwiseConv2D	(1, 1, 128)	1,152	3,600 (1x1x128x3x3x1)
Conv2D	(1, 1, 128)	16,512	128 (1x1x128x1x1x128)
BatchNormalization	(1, 1, 128)	512	0
GlobalAveragePooling2D	(None, 128)	0	0
Flatten	(None, 128)	0	0
Dense	(None, 128)	16,512	16,512
BatchNormalization	(None, 128)	512	0
Dense	(None, 10)	1,290	1,280

**Did you observe any overfitting? Should the model train for longer, shorter, or about that number of epochs.**

The validation accuracy is lower than the training accuracy. The validation loss is higher than the training loss. So, this model is overfitting.



**Does it correctly label the picture?**

Sample image of Dog is tested on Model 2 and it correctly classified.

```

img = Image.open("dog.jpg") # Loading Image
img = img.resize((32, 32)) # Resizing
x = np.array(img) # Convert the image to a numpy array
x = np.expand_dims(x, axis=0)
model = keras.models.load_model("model2.h5")
# Make a prediction on the image
pred = model.predict(x)
# Get the predicted class label
class_index= np.argmax(pred)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
class_name = class_names[class_index]
print(class_name)

1/1 [=====] - 0s 196ms/step
dog

```



### Model3

Layer	Output Shape	Parameters	MAC Operations
InputLayer	(None, 32, 32, 3)	0	0
Conv2D	(None, 16, 16, 32)	896	294,912 (16x16x32x3x3x3x32)
BatchNormalization	(None, 16, 16, 32)	128	0
Dropout	(None, 16, 16, 32)	0	0
Conv2D	(None, 8, 8, 64)	18,496	36,864 (8x8x64x3x3x32)
BatchNormalization	(None, 8, 8, 64)	256	0
Dropout	(None, 8, 8, 64)	0	0
Conv2D	(None, 4, 4, 128)	73,856	36,864 (4x4x128x3x3x64)
BatchNormalization	(None, 4, 4, 128)	512	0
Dropout	(None, 4, 4, 128)	0	0
Conv2D	(None, 4, 4, 128)	4,224	11,520 (4x4x128x1x1x32)
Add	(None, 4, 4, 128)	0	0
Dropout	(None, 4, 4, 128)	0	0
Conv2D	(None, 2, 2, 128)	1,47,584	11,520 (2x2x128x3x3x1)
BatchNormalization	(None, 2, 2, 128)	512	0
Dropout	(None, 2, 2, 128)	0	0
Conv2D	(None, 1, 1, 128)	1,47,584	32 (1x1x128x1x1x128)
BatchNormalization	(None, 1, 1, 128)	512	0
Dropout	(None, 1, 1, 128)	0	0
Add	(None, 4, 4, 128)	0	0
Dropout	(None, 4, 4, 128)	0	0
Conv2D	(None, 4, 4, 128)	1,47,584	36,864 (4x4x128x3x3x32)
BatchNormalization	(None, 4, 4, 128)	512	0
Dropout	(None, 4, 4, 128)	0	0
Conv2D	(None, 4, 4, 128)	1,47,584	36,864 (4x4x128x3x3x32)
BatchNormalization	(None, 4, 4, 128)	512	0
Dropout	(None, 4, 4, 128)	0	0
Add	(None, 4, 4, 128)	0	0
Dropout	(None, 4, 4, 128)	0	0
MaxPooling2D	(None, 1, 1, 128)	0	0

Flatten	(None, 128)	0	0
Dense	(None, 128)	16,512	16,512
BatchNormalization	(None, 128)	512	0
Dense	(None, 10)	1,290	1,280

**Did you observe any overfitting? Should the model train for longer, shorter, or about that number of epochs.**

The validation accuracy is almost similar to the training accuracy. The validation loss is nearly same as training loss. So, this model is slightly overfitting.



## Does it correctly label the picture?

A picture of an dog was tested, and Model3 Classified it as Truck.

```
img = Image.open("dog.jpg") # Loading Image
img = img.resize((32, 32)) # Resizing
x = np.array(img) # Convert the image to a numpy array
x = np.expand_dims(x, axis=0)
model = keras.models.load_model("model3.h5")
# Make a prediction on the image
pred = model.predict(x)
# Get the predicted class label
class_index= np.argmax(pred)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
class_name = class_names[class_index]
print(class_name)
```

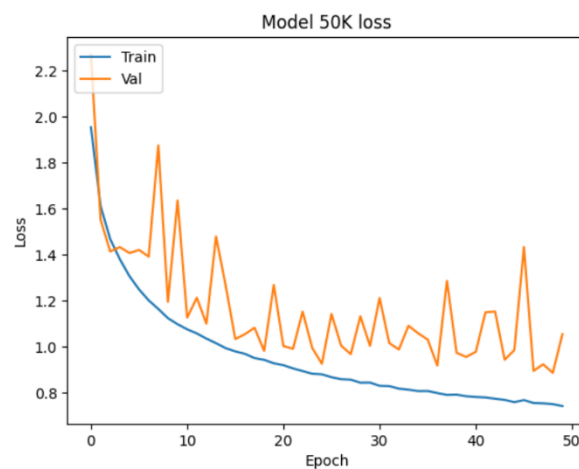

```
WARNING:tensorflow:6 out of the last 6 calls to <function Model.make_predict_function.<locals>.predict_function
1/1 [=====] - 0s 209ms/step
truck
```

## Model\_50K

Layer	Output Shape	Parameters	MAC Operations
Conv2D	(None, 16, 16, 32)	416	819,200 (16x16x32x3x3x32)
BatchNormalization	(None, 16, 16, 32)	128	0
Activation	(None, 16, 16, 32)	0	0
SeparableConv2D	(None, 8, 8, 64)	2,400	36,864 (8x8x64x3x3x32)
BatchNormalization	(None, 8, 8, 64)	256	0
SeparableConv2D	(None, 4, 4, 128)	8,576	36,864 (4x4x128x3x3x64)
BatchNormalization	(None, 4, 4, 128)	512	0
Dropout	(None, 4, 4, 128)	0	0
SeparableConv2D	(None, 4, 4, 128)	17,664	36,864 (4x4x128x3x3x128)
BatchNormalization	(None, 4, 4, 128)	512	0
Dropout	(None, 4, 4, 128)	0	0
MaxPooling2D	(None, 1, 1, 128)	0	0
Flatten	(None, 128)	0	0
Dense	(None, 128)	16,512	16,512
Dropout	(None, 128)	0	0
Dense	(None, 10)	1,290	1,280

**Did you observe any overfitting? Should the model train for longer, shorter, or about that number of epochs.**

The validation accuracy is similar to the training accuracy. The validation loss is also the same as training loss. Model\_50K performing better than previous models



### Does it correctly label the picture?

A picture of a dog was tested, and it classified it as cat.

```
img = Image.open("dog.jpg") # Loading Image
img = img.resize((32, 32)) # Resizing
x = np.array(img) # Convert the image to a numpy array
x = np.expand_dims(x, axis=0)
model = keras.models.load_model("best_model.h5")
# Make a prediction on the image
pred = model.predict(x)
# Get the predicted class label
class_index= np.argmax(pred)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
class_name = class_names[class_index]
print(class_name)
```

```
1/1 [=====] - 0s 374ms/step
cat
```

**How did the three models compare? Consider final accuracy, time per epoch to train, number of epochs needed to reach a given accuracy, overfitting.**



The model50K exhibited the highest accuracy, reaching 0.742. The average training time per epoch ranged from 4 to 5 seconds. Training the models for 50 epochs proved effective, with overfitting showing improvement from model 1 to model50K.

**What did you observe? What model architecture or training decisions made the most difference?**

Reducing the number of model layers has had an impact on accuracy. As model layers are reduced, accuracy increases and overfitting decreases. Additionally, this reduction decreases the number of parameters.